



DLCV Homework 3

Problem 1: GAN


1. Describe the architecture & implementation details of your model.



Generator :



Discriminator :



```
class Generator(nn.Module):
    def __init__(self, ngpu, figsize=64):
        super(Generator, self).__init__()
        self.ngpu = ngpu

        self.transconv1 = nn.ConvTranspose2d(101, figsize * 8, 4, 1, 0, bias=False)
        self.batch1 = nn.BatchNorm2d(figsize * 8)
        self.relu1 = nn.ReLU()

        self.transconv2 = nn.ConvTranspose2d(figsize * 8, figsize * 4, 4, 2, 0, bias=False)
        self.batch2 = nn.BatchNorm2d(figsize * 4)
        self.relu2 = nn.ReLU()

        self.transconv3 = nn.ConvTranspose2d(figsize * 4, figsize * 2, 4, 2, 0, bias=False)
        self.batch3 = nn.BatchNorm2d(figsize * 2)
        self.relu3 = nn.ReLU()

        self.transconv4 = nn.ConvTranspose2d(figsize * 2, figsize * 1, 4, 2, 1, bias=False)
        self.batch4 = nn.BatchNorm2d(figsize * 1)
        self.relu4 = nn.ReLU()

        self.transconv5 = nn.ConvTranspose2d(figsize * 1, figsize * 1, 3, 4, 2, 1, bias=False)
        self.tanh = nn.Tanh()
```

```
class Discriminator(nn.Module):
    def __init__(self, ngpu, figsize=64):
        super(Discriminator, self).__init__()
        self.ngpu = ngpu

        self.conv1 = nn.Conv2d(3, figsize * 4, 2, 1, 0, bias=False)
        self.leakyrelu1 = nn.LeakyReLU(0.2, inplace=True)

        self.conv2 = nn.Conv2d(figsize * 4, figsize * 2, 4, 2, 1, bias=False)
        self.batch2 = nn.BatchNorm2d(figsize * 2)
        self.leakyrelu2 = nn.LeakyReLU(0.2, inplace=True)

        self.conv3 = nn.Conv2d(figsize * 2, figsize * 1, 4, 2, 1, bias=False)
        self.batch3 = nn.BatchNorm2d(figsize * 1)
        self.leakyrelu3 = nn.LeakyReLU(0.2, inplace=True)

        self.conv4 = nn.Conv2d(figsize * 1, figsize * 1, 4, 2, 1, bias=False)
        self.batch4 = nn.BatchNorm2d(figsize * 1)
        self.leakyrelu4 = nn.LeakyReLU(0.2, inplace=True)

        self.conv5 = nn.Conv2d(figsize * 1, figsize * 1, 1, 4, 1, 0, bias=False)
        self.sigmoid = nn.Sigmoid()
```

Implementation details :

batch size = 128	loss function = BCEloss
latent dimension = 101	optimizers = AdamOptimizer for both generator and discriminator
learning rate = 0.0002	normalized images with : mean = 0.5 std = 0.5
beta1 = 0.5	latent vectors = from normal distribution (mean = 0 , std = 1)
number of epochs = 130	

2. Plot 32 random images generated from your model. [fig1_2.jpg]



fig1_2.jpg

3. Discuss what you've observed and learned from implementing GAN.

The discriminator loss drops a lot more faster and to a lower score than the generator : it can indeed drop under 0.02 while the generator loss is still around 0.1. But when the discriminator loss is very low the generator doesn't produce as good quality image as when the two losses are more even.


Also I noted that for two models with roughly the same generator loss but different discriminator loss, the lower the discriminator loss the better the output images whereas for two models with the same discriminator loss but different generator losses there wasn't a significant difference of quality between outputs. Thus we can think that the discriminator impacts the quality more than the generator.

Another point is that the quality of the generated faces is also dependent on the randomly created input noise.

1. Describe the architecture & implementation details of your model.

Diagram illustrating the components of a neural network architecture, showing 10 types of layers and their corresponding 3D block representations:

- transpose convolution layer (blue block)
- batch normalization layer (yellow block)
- Relu layer (red block)
- tanh layer (green block)
- convolution layer (light blue block)
- LeakyRelu layer (orange block)
- sigmoid layer (purple block)
- softmax layer (olive block)
- linear layer (grey block)



Output shape : 1024 x 4 x 4 512 x 8 x 8 256 x 16 x 16 128 x 32 x 32 3 x 64 x 64

Output shape : 128 x 32 x 32 256 x 16 x 16 512 x 8 x 8 1024 x 4 x 4 64 x 1 x 1 1 x 1 x 1

```
class Discriminator(nn.Module):
    def __init__(self, figsize=64):
        super(Discriminator, self).__init__()
        self.decoder = nn.Sequential(
            nn.Conv2d(3, figsize, 4, 2, 1),
            nn.LeakyReLU(0.2, inplace=True),
            nn.Conv2d(figsize, figsize * 2, 4, 2, 1),
            nn.BatchNorm2d(figsize * 2),
            nn.LeakyReLU(0.2, inplace=True),
            nn.Conv2d(figsize * 2, figsize * 4, 4, 2, 1),
            nn.BatchNorm2d(figsize * 4),
            nn.LeakyReLU(0.2, inplace=True),
            nn.Conv2d(figsize * 4, figsize * 8, 4, 2, 1),
            nn.BatchNorm2d(figsize * 8),
            nn.LeakyReLU(0.2, inplace=True),
            nn.Conv2d(figsize * 8, figsize * 1, 4, 1, 0),
        )
        self.fc_dis = nn.Linear(figsize * 1, 1)
        self.fc_aux = nn.Linear(figsize * 1, 1)

        self.softmax = nn.Softmax()
        self.sigmoid = nn.Sigmoid()
```

Implementation details :

batch size = 64	loss function = BCEloss
latent dimension = 101	optimizers = AdamOptimizer for both generator and discriminator
learning rate = 0.0002	normalized images with : mean = 0.5 std = 0.5
beta1 = 0.5	latent vectors = from normal distribution (mean = 0 , std = 1)
number of epochs = 200	

2. Plot 10 random pairs of generated images from your model, where each pair should be generated from the same random vector input but with opposite attribute. This is to demonstrate your model's ability to disentangle features of interest. [fig2_2.jpg]



fig2_2.jpg

3. Discuss what you've observed and learned from implementing ACGAN.

The generator of ACGAN is the same as the generator of problem1 GAN and the discriminators are also almost the same except than in ACGAN it returns two informations (data and attribute) but the quality of the output images is a lot more greater in this problem which I can only explain by the fact that the training of ACGAN lasted five times longer than the training of GAN.

It is also noticeable that the mouth on which is applied (or not) the "smiling" attribute is often the part of the face that is the most well-designed and realistic. The generator therefore might perform better with constraint because it's not only constraint but also further information to generate from.

Problem 3: DANN


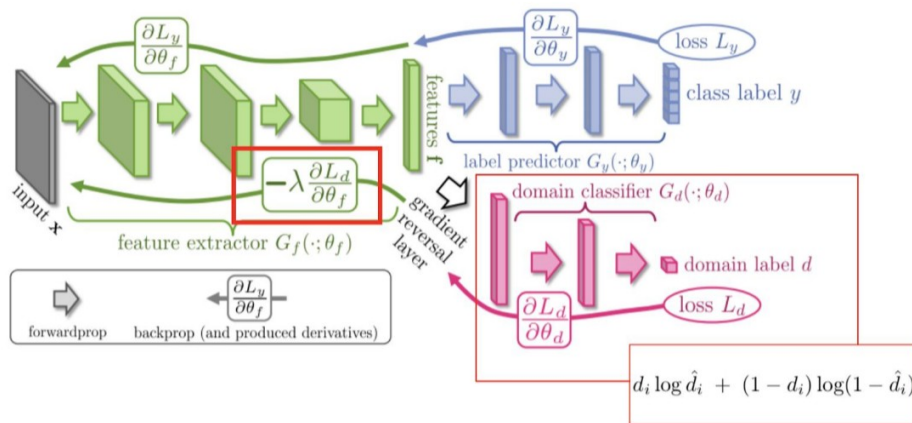
1. Compute the accuracy on target domain, while the model is trained on source domain only. (lower bound)
2. Compute the accuracy on target domain, while the model is trained on source and target domain. (domain adaptation)
3. Compute the accuracy on target domain, while the model is trained on target domain only. (upper bound)

	SVHN \rightarrow MNIST-M	MNIST-M \rightarrow SVHN
Trained on source	MNIST dataset: 0.328200	SVHN dataset: 0.375576
Adaptation (DANN/Improved)	DANN : 89.14% UDA : 89.87%	DANN : 41.057160417947145% UDA : 42.689766441303014%
Trained on target	MNIST dataset: 0.947100	SVHN dataset: 0.740934


4. Visualize the latent space by mapping the testing images to 2D space (with t-SNE) and use different colors to indicate data of (a) different digit classes 0-9 and (b) different domains (source/target).

Not done.

5. Describe the architecture & implementation detail of your model.



Feature extractor :




Output shape :

64 x 13 x 13

128 x 5 x 5

Label predictor :




Output shape :

100

10

Domain classifier :



Output shape :

2

`class DANN_Neural_Network(nn.Module):`

```
def __init__(self):
    super(DANN_Neural_Network, self).__init__()
    self.attr = nn.Sequential(
        nn.Conv2d(3, 64, kernel_size=5),
        nn.BatchNorm2d(64),
        nn.MaxPool2d(2),
        nn.ReLU(True),
        nn.Conv2d(64, 50, kernel_size=5),
        nn.BatchNorm2d(50),
        nn.Dropout2d(),
        nn.MaxPool2d(2),
        nn.ReLU(True))
```

```
self.cclass = nn.Sequential(
    nn.Linear(50 * 4 * 4, 100),
    nn.BatchNorm1d(100),
    nn.ReLU(True),
    nn.Dropout2d(),
    nn.Linear(100, 100),
    nn.BatchNorm1d(100),
    nn.ReLU(True),
    nn.Linear(100, 10),
    nn.LogSoftmax())
```

```
self.domain = nn.Sequential(
    nn.Linear(50 * 4 * 4, 100),
    nn.BatchNorm1d(100),
    nn.ReLU(True),
    nn.Linear(100, 2),
    nn.LogSoftmax(dim=1))
```

Implementation details :

batch size = 128 loss function = NLLLoss
learning rate = 0.0001 normalized images with : mean = 0.5 std = 0.5
beta1 = 0.5 number of epochs = 100
optimizer = AdamOptimizer

6. Discuss what you've observed and learned from implementing DANN.

When the loss of the domain classifier was more than two times the loss of the feature extractor it made the loss of the class classifier rise: so when the domain classifier was outperformed by the feature extractor it then led to wrong label classifications. Therefore the capacity of the feature extractor, class classifier and domain classifier should be around the same level during the training.

The accuracy is very different whether we are in (1) svhn to mnist or (2) mnist to svhn. I think it may be because the ratio between test and training are also very different whether we take the mnist dataset (1/6) or the svhn dataset (1/4) so that models are relatively more trained on mnist. It's especially true for the adaptation model which performs a lot more on the mnist test dataset.

Problem 4: Improved UDA model

I implemented the ADDA model.

1. Compute the accuracy on target domain, while the model is trained on source and target domain. (domain adaptation)

On SVHN test dataset :

42.689766441303014%

On MNIST test dataset :


89.87%

2. Visualize the the latent space by mapping the testing images to 2D space (with t-SNE) and use different colors to indicate data of (a) different digits classes 0-9 and (b) different domains (source/target).

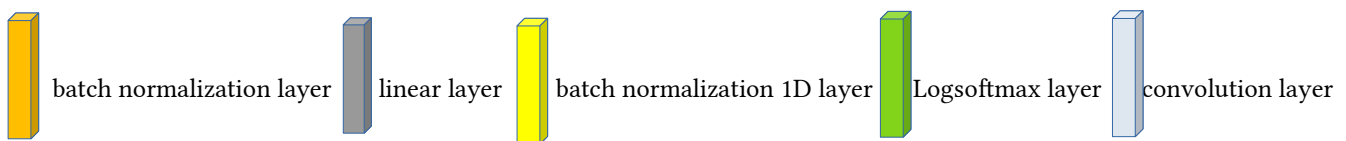
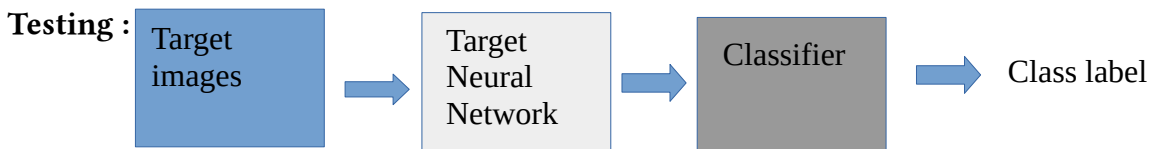
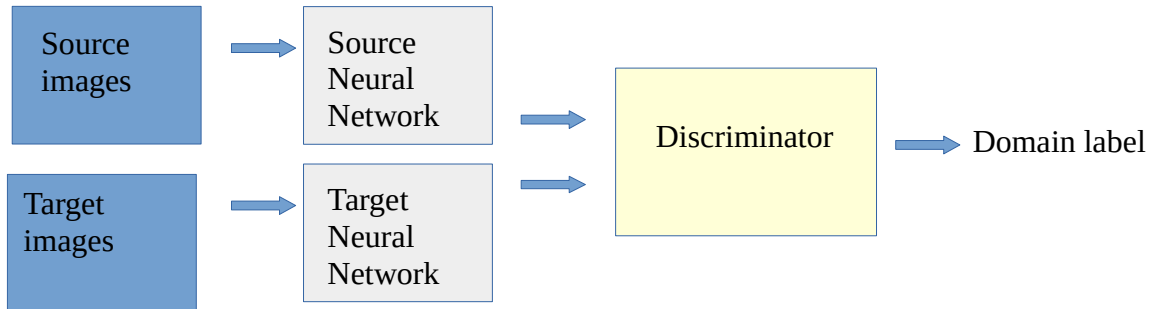
Not done.

3. Describe the architecture & implementation detail of your model.


ADDA concept:



Adversarial Adaptation :




Neural Network (source and target)



Output shape :

500


Classifier



Output shape :

10

Discriminator



Output shape :

2

```

class UDA_NeuralNetwork(nn.Module):
    def __init__(self):
        super(UDA_NeuralNetwork, self).__init__()
        self.conv1 = nn.Conv2d(1, 20, 5)
        self.bn1 = nn.BatchNorm2d(20)
        self.conv2 = nn.Conv2d(20, 50, 5)
        self.bn2 = nn.BatchNorm2d(50)
        self.fc1 = nn.Linear(50 * 4 * 4, 500)
        self.bn3 = nn.BatchNorm1d(500)
  
```

```

class UDA_NeuralNetwork_Classifier(nn.Module):
    def __init__(self):
        super(UDA_NeuralNetwork_Classifier, self).__init__()
        self.fc2 = nn.Linear(500, 10)
        self.lsftm = nn.Logsoftmax()
  
```

```

class Discriminator(nn.Module):
    def __init__(self):
        super(Discriminator, self).__init__()
        self.fc1 = nn.Linear(500, numberHiddenUnitsD)
        self.fc2 = nn.Linear(numberHiddenUnitsD, numberHiddenUnitsD)
        self.fc3 = nn.Linear(numberHiddenUnitsD, 2)
        self.bn1 = nn.BatchNorm1d(numberHiddenUnitsD)
        self.bn2 = nn.BatchNorm1d(numberHiddenUnitsD)
  
```


Implementation details :

batch size = 256	loss function = CrossEntropyLoss
learning rate = 0.0002	normalized images with : mean = 0.5 std = 0.5
beta1 = 0.5	optimizer = AdamOptimizer for discriminator and targetNN

4. Discuss what you've observed and learned from implementing your improved UDA model.

The ADDA model was more difficult to train than the DANN model, it's more likely due to the fact that there is two different feature extractors involved in the process but the accuracy results aren't significantly much better than the results of DANN considering the training was a lot longer.

Also we can make the same remarks as for the DANN model in regards to the difference of performance on the mnist test dataset and on the svhn test dataset.

References :

https://pytorch.org/tutorials/beginner/dcgan_faces_tutorial.html

<https://stephan-osterburg.gitbook.io/coding/coding/ml-dl/tensorfow/chapter-4-conditional-generative-adversarial-network/implementation-of-acgan-model>

<https://github.com/corenel/pytorch-adda/tree/master/models>

No collaborator